


The Price is Right

Helping Homeowners Find The Right Price to Sell Their Home

```
In [1]: from IPython.display import Image  
Image("/Users/matthewnykaza/Documents/Flatiron/Phase_2_Project/dsc-phase-2-
```

Out[1]:



Overview

For this project I was given, by the Flatiron school, a dataset consisting of the King Country Housing sales over a period of May 2014 to May 2015. I was tasked with finding a relevant business problem and conducting analysis using the given data and conducting statistical models to answer the business question. I decided to try and help basic homeowners find the appropriate price to list their home at, given basic variable (bedrooms, bathrooms, lot size etc.). I began this process by preprocessing the data by taking out unnecessary variables (i.e. views, sqft basement, etc.) and changing some of the variables (changing basement size to whether or not the home has a basement for example). Once this basic data cleaning was complete I set out to create a base model just to see what the correlations, r2 score and other relevant testing metrics looked like before doing any standardization techniques. With that subpar model complete I began to change the data using standardization and normalizing techniques. This started with log-transforming the target variable (selling price) which turned that variable into a much more normal distribution. Following that I used a standard scaler from sklearn which removed the mean from the

independant variables and scaled to the variance. This did not necessarily make the data "more normal", but it did bring center the data around a mean of 0 with a standard deviation of 1. This did improve on my model somewhat, but overall did not make my prediction power (as read by R2 score) any stronger. My final model featured log-transformed, and Standard Scaled data, which did not improve the R2 score, but did provide the least error, and gives me confidence that there is something strong to go on here. I recommend that a home seller not use this model quite yet, but with more time I am confident that there can be a great conclusion reached, and a accurate home selling model.

Business Problem

Selling a home can be an exceedingly stressful process. One of the biggest pain points for home sellers is selecting the correct listing price for that home. Pick a price that is too low, and you run the risk of losing out on a lot of potential earnings. To high and you can wind up with a house that remains on the market for a long time, potentially hamstringing efforts to move into that new house. With this dataset, and statistical knowledge I intend to create a model that will assist with people knowing the right price to list their home at.

```
In [72]: #Start with loading in all the libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.preprocessing import PolynomialFeatures
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import RobustScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import Lasso, Ridge, LinearRegression
from sklearn.metrics import r2_score, mean_squared_error, make_scorer
from sklearn.feature_selection import RFE
from scipy import stats
import statsmodels.api as sm
import statsmodels.formula.api as smf
```

King Count Housing Dataset

- In this modeling I used data provided to me by Flatiron School. It consists of a single dataset that include relevant information about housing sales from May 2014 to May 2015 in King County, WA.
- There are many variables in this dataset and they include home ID (a simple identifier unique to each sold home), selling price, number of bedrooms and bathrooms, square feet of the living space, squarefeet of the total lot, number of floors, whether the home is on the waterfront or not, whether the home has been viewed or not, a overall condition rating, a King County

Housing grade, year built, year remodeled, zipcode it is in, the latitude and longitude of the location, and the square feet of the 15 surrounding homes lots and living space.

- The target variable for this modeling will be the selling price ('price' in the data).
- The feature variables that I will be using for this project will be bedrooms, bathrooms, sqft_lot, floors, waterfront, condition, grade, and year built.

```
In [73]: #Load in the dataset and take a look at its overview
data = pd.read_csv('data/kc_house_data.csv')
```

```
In [74]: #Take a look at the first 5 rows of data
data.head()
```

Out[74]:

| | id | date | price | bedrooms | bathrooms | sqft_living | sqft_lot | floors | waterfront |
|---|------------|------------|----------|----------|-----------|-------------|----------|--------|------------|
| 0 | 7129300520 | 10/13/2014 | 221900.0 | 3 | 1.00 | 1180 | 5650 | 1.0 | NaN |
| 1 | 6414100192 | 12/9/2014 | 538000.0 | 3 | 2.25 | 2570 | 7242 | 2.0 | 0.0 |
| 2 | 5631500400 | 2/25/2015 | 180000.0 | 2 | 1.00 | 770 | 10000 | 1.0 | 0.0 |
| 3 | 2487200875 | 12/9/2014 | 604000.0 | 4 | 3.00 | 1960 | 5000 | 1.0 | 0.0 |
| 4 | 1954400510 | 2/18/2015 | 510000.0 | 3 | 2.00 | 1680 | 8080 | 1.0 | 0.0 |

5 rows × 21 columns

```
In [75]: #Check out the columns
data.columns
```

```
Out[75]: Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',
               'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade',
               'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode',
               'lat', 'long', 'sqft_living15', 'sqft_lot15'],
              dtype='object')
```

```
In [76]: #Check out the basic stats for the dataset
data.describe()
```

Out[76]:

| | id | price | bedrooms | bathrooms | sqft_living | sqft_lot | |
|--------------|--------------|--------------|--------------|--------------|--------------|--------------|-------|
| count | 2.159700e+04 | 2.159700e+04 | 21597.000000 | 21597.000000 | 21597.000000 | 2.159700e+04 | 21597 |
| mean | 4.580474e+09 | 5.402966e+05 | 3.373200 | 2.115826 | 2080.321850 | 1.509941e+04 | 1. |
| std | 2.876736e+09 | 3.673681e+05 | 0.926299 | 0.768984 | 918.106125 | 4.141264e+04 | 0 |
| min | 1.000102e+06 | 7.800000e+04 | 1.000000 | 0.500000 | 370.000000 | 5.200000e+02 | 1. |
| 25% | 2.123049e+09 | 3.220000e+05 | 3.000000 | 1.750000 | 1430.000000 | 5.040000e+03 | 1. |
| 50% | 3.904930e+09 | 4.500000e+05 | 3.000000 | 2.250000 | 1910.000000 | 7.618000e+03 | 1. |
| 75% | 7.308900e+09 | 6.450000e+05 | 4.000000 | 2.500000 | 2550.000000 | 1.068500e+04 | 2 |
| max | 9.900000e+09 | 7.700000e+06 | 33.000000 | 8.000000 | 13540.000000 | 1.651359e+06 | 3 |

A few things noticed here

1. Outliers in a bunch of the categories
 - Especially bedrooms, bathrooms, sqft_lot and sqft_lot15

```
In [77]: #Look for any NaN data
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   id                    21597 non-null  int64
1   date                 21597 non-null  object
2   price               21597 non-null  float64
3   bedrooms            21597 non-null  int64
4   bathrooms           21597 non-null  float64
5   sqft_living         21597 non-null  int64
6   sqft_lot            21597 non-null  int64
7   floors              21597 non-null  float64
8   waterfront          19221 non-null  float64
9   view                21534 non-null  float64
10  condition            21597 non-null  int64
11  grade               21597 non-null  int64
12  sqft_above          21597 non-null  int64
13  sqft_basement       21597 non-null  object
14  yr_built            21597 non-null  int64
15  yr_renovated        17755 non-null  float64
16  zipcode             21597 non-null  int64
17  lat                 21597 non-null  float64
18  long                21597 non-null  float64
19  sqft_living15       21597 non-null  int64
20  sqft_lot15          21597 non-null  int64
dtypes: float64(8), int64(11), object(2)
memory usage: 3.5+ MB
```

A few things noticed here

1. We have a few NaNs
2. date and swft_basement are the only two columns with a data type as "object", this makes sense for date, but why sqft_basement?

```
In [78]: #As show in data.info() sqft_basement is an object, but why
print(data.groupby('sqft_basement')['id'].nunique())
```

```
sqft_basement
0.0          12718
10.0           1
100.0          42
1000.0         146
1008.0          1
...
960.0          65
970.0          44
980.0          55
990.0          51
?             454
Name: id, Length: 304, dtype: int64
```

Some 454 are listed as ?, we're just going to make an assumption that this means they don't have a basement because in my opinion it is pretty easy for someone to tell if a place does or not, we will make these values 0.

```
In [79]: #Change '?' to 0 in sqft_basement and check
data['sqft_basement'] = data['sqft_basement'].replace(['?'], 0)
```

```
In [80]: #Sanity check
data.loc[lambda df: df['sqft_basement'] == '?']
```

```
Out[80]:
```

| id | date | price | bedrooms | bathrooms | sqft_living | sqft_lot | floors | waterfront | view | ... | grade | sqft_basement |
|---------------------|------|-------|----------|-----------|-------------|----------|--------|------------|------|-----|-------|---------------|
| 0 rows × 21 columns | | | | | | | | | | | | |

```
In [81]: #Now change them to numeric in the same way the other sqft variables are
data['sqft_basement'] = data[['sqft_basement']].apply(pd.to_numeric)
```

```
In [82]: #Sanity check
data['sqft_basement'].dtype
```

```
Out[82]: dtype('float64')
```

```

In [83]: # Idea on distribution of data with a histogram
#Going to create a function that will make this process easier in the future
#times throughout the process
#Lets create a histogram to check for normality
def histogram_view(data_set):
    """This will produce 3 columns of histograms and a corresponding number
    of rows depending on the len(data_set.columns). The result will be a
    side by side view of all of the histograms of each column in a data set
    ncols = 3
    nrows = int(np.ceil(len(data_set.columns) / (1.0*ncols)))
    fig, axes = plt.subplots(nrows=nrows, ncols=ncols, figsize=(20, 15))
    # Lazy counter so we can remove unwanted axes
    counter = 0
    for i in range(nrows):
        for j in range(ncols):

            ax = axes[i][j]

            # Plot when we have data
            if counter < len(data_set.columns):

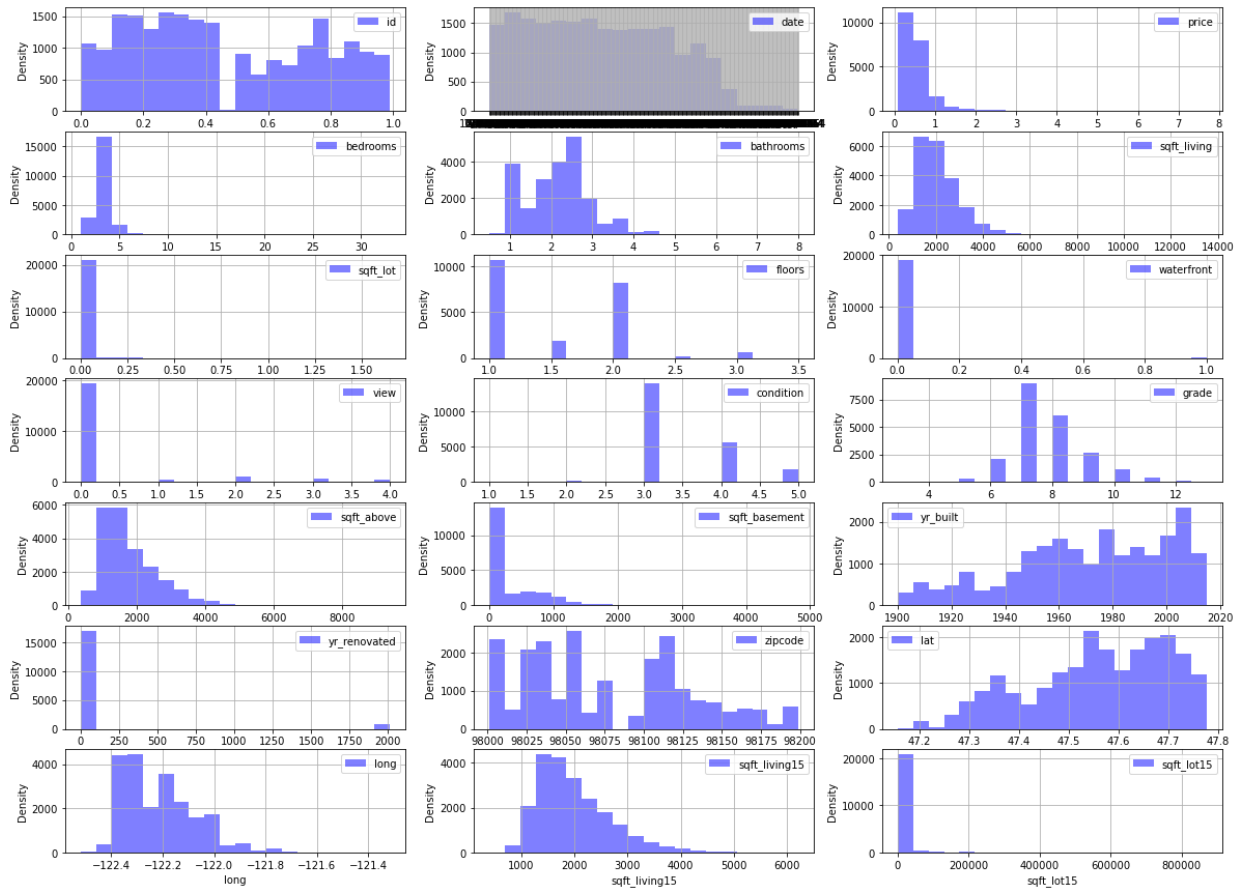
                ax.hist(data_set[data_set.columns[counter]], bins=20, color='red')
                ax.set_xlabel('{}'.format(data_set.columns[counter]))
                ax.set_ylabel('Density')
                leg = ax.legend(loc='best')
                leg.draw_frame(True)
                ax.grid(which='both', axis='both', linestyle='--')
            # Remove axis when we no longer have data
            else:
                ax.set_axis_off()

            counter += 1

plt.show()

```

```
In [84]: histogram_view(data)
```



We can already get the sense that much of the data is far from normal especially in the cases of id, date, and sqft_basement. For now we are not worrying about the normality of the categorical variables which look to be bedrooms, bathrooms, floors, waterfront, view, condition, grade, yr_built, yr_renovated, zipcode, lat and long.

Data preprocessing

- First I just wanted to make a basic check for any duplicates in the data
- NaN data was addressed and major outliers were identified and corrected accordingly
- Upon starting this project I decided to drop the date, view, sqft_above, yr_renovated, zipcode, lat, long, sqft_living15 and sqft_lot15 the reasoning for that can be seen below.
- After dropping those named variables I went to check about any NaN data.
 - The only category was the Waterfront category, and handling of that issue can be seen below.
- Then I took a look for outliers which

Check for Duplicates

- Will be using the ID column for this

```
In [85]: duplicates = data[data.id.duplicated(keep=False)]
duplicates
```

Out[85]:

| | id | date | price | bedrooms | bathrooms | sqft_living | sqft_lot | floors | water |
|--------------|------------|------------|-----------|----------|-----------|-------------|----------|--------|-------|
| 93 | 6021501535 | 7/25/2014 | 430000.0 | 3 | 1.50 | 1580 | 5000 | 1.0 | |
| 94 | 6021501535 | 12/23/2014 | 700000.0 | 3 | 1.50 | 1580 | 5000 | 1.0 | |
| 313 | 4139480200 | 6/18/2014 | 1380000.0 | 4 | 3.25 | 4290 | 12103 | 1.0 | |
| 314 | 4139480200 | 12/9/2014 | 1400000.0 | 4 | 3.25 | 4290 | 12103 | 1.0 | |
| 324 | 7520000520 | 9/5/2014 | 232000.0 | 2 | 1.00 | 1240 | 12092 | 1.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 20654 | 8564860270 | 3/30/2015 | 502000.0 | 4 | 2.50 | 2680 | 5539 | 2.0 | |
| 20763 | 6300000226 | 6/26/2014 | 240000.0 | 4 | 1.00 | 1200 | 2171 | 1.5 | |
| 20764 | 6300000226 | 5/4/2015 | 380000.0 | 4 | 1.00 | 1200 | 2171 | 1.5 | |
| 21564 | 7853420110 | 10/3/2014 | 594866.0 | 3 | 3.00 | 2780 | 6000 | 2.0 | |
| 21565 | 7853420110 | 5/4/2015 | 625000.0 | 3 | 3.00 | 2780 | 6000 | 2.0 | |

353 rows × 21 columns

There looks to be some duplicates, but all of them seem to make sense and simply looks like they are the same home being resold numerous times. With the duplicates only representing roughly 2% of the data we are not going to worry about them.

NaN Data

```
In [86]: #Dealing with NaN's
data.info()
#Looks to be only the waterfront column, so lets dig into that one
```

| | | | | |
|----|---------------|-------|----------|---------|
| 3 | bedrooms | 21597 | non-null | int64 |
| 4 | bathrooms | 21597 | non-null | float64 |
| 5 | sqft_living | 21597 | non-null | int64 |
| 6 | sqft_lot | 21597 | non-null | int64 |
| 7 | floors | 21597 | non-null | float64 |
| 8 | waterfront | 19221 | non-null | float64 |
| 9 | view | 21534 | non-null | float64 |
| 10 | condition | 21597 | non-null | int64 |
| 11 | grade | 21597 | non-null | int64 |
| 12 | sqft_above | 21597 | non-null | int64 |
| 13 | sqft_basement | 21597 | non-null | float64 |
| 14 | yr_built | 21597 | non-null | int64 |
| 15 | yr_renovated | 17755 | non-null | float64 |
| 16 | zipcode | 21597 | non-null | int64 |
| 17 | lat | 21597 | non-null | float64 |
| 18 | long | 21597 | non-null | float64 |
| 19 | sqft_living15 | 21597 | non-null | int64 |
| 20 | sqft_lot15 | 21597 | non-null | int64 |

```
dtypes: float64(9), int64(11), object(1)
memory usage: 3.5+ MB
```

```
In [87]: #Let's do a quick check on the counts of the waterfront column to give us a
print(data.groupby('waterfront')['id'].nunique())
```

```
waterfront
0.0    18941
1.0     146
Name: id, dtype: int64
```

As determined by the histogram the waterfront variable is a categorical variable. There are roughly 2,000 NaNs in that variable, but with only 146 waterfront homes in total on the dataset I believe that instead of removing the NaNs it is pretty safe to assume that these NaNs are non-waterfront homes and will simply fill them in at 0 indicating they are not on a waterfront.

In future modeling this can be improved by comparing the NaNs and their lat and long values to verify whether or not this is actually the case.

```
In [88]: #Fill with 0 and sanity check
data['waterfront'] = data['waterfront'].fillna(0)
data.info()
```

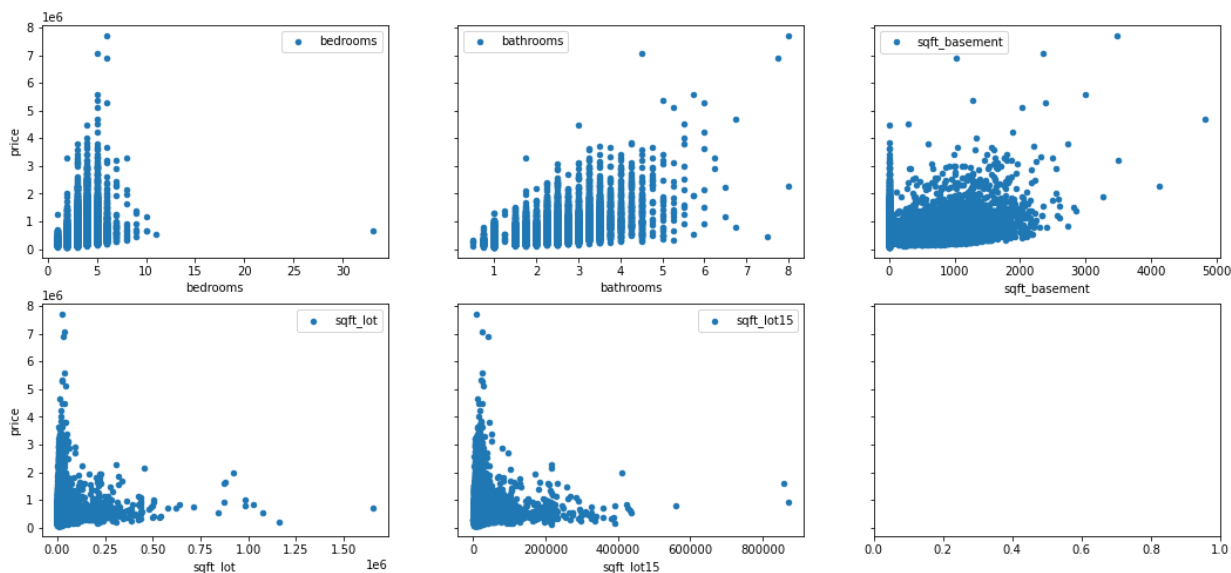
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21597 non-null  int64
1   date                   21597 non-null  object
2   price                  21597 non-null  float64
3   bedrooms               21597 non-null  int64
4   bathrooms              21597 non-null  float64
5   sqft_living            21597 non-null  int64
6   sqft_lot               21597 non-null  int64
7   floors                 21597 non-null  float64
8   waterfront             21597 non-null  float64
9   view                   21534 non-null  float64
10  condition              21597 non-null  int64
11  grade                  21597 non-null  int64
12  sqft_above             21597 non-null  int64
13  sqft_basement          21597 non-null  float64
14  yr_built                21597 non-null  int64
```

We have confirmed that all outliers were successfully removed from waterfront, now we do have some in the view and yr_renovated categories, but we will not worry about those will be removed, but more on that later.

Outliers

We're just going to look at the previously highlighted columns of bedrooms, bathrooms, sqft_lot and sqft_lot15

```
In [89]: outlier_col = ['bedrooms', 'bathrooms', 'sqft_basement', 'sqft_lot', 'sqft_
fig, axes = plt.subplots(2,3,sharey=True, figsize=(18,8))
axe = axes.ravel()
for idx, column in enumerate(data[outlier_col]):
    data.plot(kind='scatter', x=column, y='price', ax=axe[idx], label=column
```



1. bedrooms has a clear outlier of 33, going to remove that one
2. bathrooms has a few outliers as well, but those are a little closer together, and I think we'll have a different method of dealing with that variable
3. sqft_basement certainly has some outliers, but they too do not look too out of whack, and we have a different plan for that variable as to be seen later
4. One major outlier here, probably going to remove that one
5. Two major outliers here, same as sqft_lot, going to remove those

A quick look shows that there are some outliers in bedrooms, bathrooms, sqft_living, and sqft_loft. I will look at the bedrooms and bathrooms columns to get an idea of what the outliers look like.

```
In [90]: #Get a closer look at the bathrooms and bedrooms columns (should be easier
print(data.groupby('bedrooms')['id'].nunique())
print(data.groupby('bathrooms')['id'].nunique())
```

bedrooms

| | |
|----|------|
| 1 | 191 |
| 2 | 2736 |
| 3 | 9731 |
| 4 | 6849 |
| 5 | 1586 |
| 6 | 265 |
| 7 | 38 |
| 8 | 13 |
| 9 | 6 |
| 10 | 3 |
| 11 | 1 |
| 33 | 1 |

Name: id, dtype: int64

bathrooms

| | |
|------|------|
| 0.50 | 4 |
| 0.75 | 70 |
| 1.00 | 3794 |
| 1.25 | 9 |
| 1.50 | 1429 |
| 1.75 | 3020 |
| 2.00 | 1913 |
| 2.25 | 2031 |
| 2.50 | 5352 |
| 2.75 | 1182 |
| 3.00 | 747 |
| 3.25 | 586 |
| 3.50 | 729 |
| 3.75 | 155 |
| 4.00 | 134 |
| 4.25 | 79 |
| 4.50 | 99 |
| 4.75 | 23 |
| 5.00 | 21 |
| 5.25 | 13 |
| 5.50 | 10 |
| 5.75 | 4 |
| 6.00 | 6 |
| 6.25 | 2 |
| 6.50 | 2 |
| 6.75 | 2 |
| 7.50 | 1 |
| 7.75 | 1 |
| 8.00 | 2 |

Name: id, dtype: int64

For the data the bathrooms are consistent from .5 to 8 while there is a pretty heavy tail, I think that using some preprocessing techniques will be the real answer here. For the bedrooms the 33 bedroom home is a massive outlier, and I will simply take that one out of the equation.

```
In [91]: #Take all the bathrooms and round them to the nearest full number, this will
#that have very little values in them.
data.loc[data['bathrooms'] <= 1.25, 'bathrooms'] = 1
data.loc[(data['bathrooms'] >=1.25) & (data['bathrooms'] <= 2.25), 'bathroom'] = 2
data.loc[(data['bathrooms'] >=2.25) & (data['bathrooms'] <= 3.25), 'bathroom'] = 3
data.loc[(data['bathrooms'] >=3.25) & (data['bathrooms'] <= 4.25), 'bathroom'] = 4
data.loc[(data['bathrooms'] >=4.25) & (data['bathrooms'] <= 5.25), 'bathroom'] = 5
data.loc[(data['bathrooms'] >=5.25) & (data['bathrooms'] <= 6.25), 'bathroom'] = 6
data.loc[(data['bathrooms'] >=6.25) & (data['bathrooms'] <= 7.25), 'bathroom'] = 7
data.loc[(data['bathrooms'] >=7.25) & (data['bathrooms'] <= 8.25), 'bathroom'] = 8
```

```
In [92]: #Sanity check
print(data.groupby('bathrooms')['id'].nunique())
```

```
bathrooms
1.0    3877
2.0    8393
3.0    7867
4.0    1097
5.0     156
6.0     22
7.0      4
8.0      4
Name: id, dtype: int64
```

```
In [93]: #Just remove that 33 bedroom outlier, that might even just be mistyped or a
#Effect the model.
data = data[data.bedrooms != 33]
data.head()
```

Out[93]:

| | id | date | price | bedrooms | bathrooms | sqft_living | sqft_lot | floors | waterfront |
|---|------------|------------|----------|----------|-----------|-------------|----------|--------|------------|
| 0 | 7129300520 | 10/13/2014 | 221900.0 | 3 | 1.0 | 1180 | 5650 | 1.0 | 0.0 |
| 1 | 6414100192 | 12/9/2014 | 538000.0 | 3 | 2.0 | 2570 | 7242 | 2.0 | 0.0 |
| 2 | 5631500400 | 2/25/2015 | 180000.0 | 2 | 1.0 | 770 | 10000 | 1.0 | 0.0 |
| 3 | 2487200875 | 12/9/2014 | 604000.0 | 4 | 3.0 | 1960 | 5000 | 1.0 | 0.0 |
| 4 | 1954400510 | 2/18/2015 | 510000.0 | 3 | 2.0 | 1680 | 8080 | 1.0 | 0.0 |

5 rows × 21 columns

Bedrooms and bathrooms outliers have been taken care of, now we'll move on to our sqft_lot and sqft_lot15

```
In [94]: data.sort_values(by=['sqft_lot'], ascending=False)
```

Out[94]:

| | id | date | price | bedrooms | bathrooms | sqft_living | sqft_lot | floors | water |
|--------------|------------|------------|-----------|----------|-----------|-------------|----------|--------|-------|
| 1717 | 1020069017 | 3/27/2015 | 700000.0 | 4 | 1.0 | 1300 | 1651359 | 1.0 | |
| 17305 | 3326079016 | 5/4/2015 | 190000.0 | 2 | 1.0 | 710 | 1164794 | 1.0 | |
| 7640 | 2623069031 | 5/21/2014 | 542500.0 | 5 | 3.0 | 3010 | 1074218 | 1.5 | |
| 7762 | 2323089009 | 1/19/2015 | 855000.0 | 4 | 4.0 | 4030 | 1024068 | 2.0 | |
| 3945 | 722069232 | 9/5/2014 | 998000.0 | 4 | 3.0 | 3770 | 982998 | 2.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 20588 | 7899800857 | 12/15/2014 | 256950.0 | 2 | 2.0 | 1070 | 635 | 2.0 | |
| 3449 | 2559950110 | 4/22/2015 | 1230000.0 | 2 | 3.0 | 2470 | 609 | 3.0 | |
| 7582 | 6371000026 | 1/22/2015 | 367500.0 | 2 | 2.0 | 1030 | 600 | 2.0 | |
| 5821 | 1773101159 | 1/7/2015 | 250000.0 | 3 | 2.0 | 1050 | 572 | 2.0 | |
| 15729 | 9828702895 | 10/22/2014 | 700000.0 | 4 | 2.0 | 2420 | 520 | 1.5 | |

21596 rows × 21 columns

Looks like that one major 5000 sqft outlier, let's remove it as to mitigate any muddiness of the data later on

```
In [95]: #remove and check for removal
data = data[data.sqft_lot != 1651359]
data.loc[lambda df: df['sqft_lot'] == 1651359]
```

Out[95]:

| | id | date | price | bedrooms | bathrooms | sqft_living | sqft_lot | floors | waterfront | view | ... | grade | sq |
|--|----|------|-------|----------|-----------|-------------|----------|--------|------------|------|-----|-------|----|
|--|----|------|-------|----------|-----------|-------------|----------|--------|------------|------|-----|-------|----|

0 rows × 21 columns

```
In [96]: #Do the same as above for sqft_lot15
data.sort_values(by=['sqft_lot15'], ascending=False)
```

Out[96]:

| | id | date | price | bedrooms | bathrooms | sqft_living | sqft_lot | floors | waterfront |
|--------------|------------|------------|-----------|----------|-----------|-------------|----------|--------|------------|
| 9705 | 225079036 | 1/7/2015 | 937500.0 | 4 | 4.0 | 5545 | 871200 | 2.0 | |
| 20436 | 1125079111 | 4/15/2015 | 1600000.0 | 4 | 6.0 | 6530 | 871200 | 2.0 | |
| 13451 | 3420069060 | 11/7/2014 | 790000.0 | 3 | 3.0 | 2640 | 432036 | 1.5 | |
| 8655 | 3226079059 | 10/19/2014 | 549950.0 | 3 | 2.0 | 2930 | 266587 | 2.0 | |
| 3797 | 1550000463 | 8/26/2014 | 637000.0 | 4 | 4.0 | 3080 | 118918 | 2.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 20891 | 8562780540 | 12/22/2014 | 325000.0 | 2 | 2.0 | 1150 | 711 | 2.0 | |
| 20999 | 8562780530 | 3/28/2015 | 338500.0 | 2 | 2.0 | 1150 | 711 | 2.0 | |
| 513 | 2827100070 | 11/5/2014 | 290000.0 | 4 | 1.0 | 1330 | 8184 | 1.5 | |
| 20733 | 2827100075 | 7/27/2014 | 286308.0 | 2 | 2.0 | 1220 | 1036 | 3.0 | |

This time we have two outliers, but they two are by a considerable margin so we will remove them from the data

```
In [97]: data = data[data.sqft_lot15 != 871200]
data = data[data.sqft_lot15 != 858132]
data.loc[lambda df: df['sqft_lot15'] > 858131]
```

Out[97]:

| | id | date | price | bedrooms | bathrooms | sqft_living | sqft_lot | floors | waterfront | view | ... | grade | sqft_lot15 |
|--|----|------|-------|----------|-----------|-------------|----------|--------|------------|------|-----|-------|------------|
|--|----|------|-------|----------|-----------|-------------|----------|--------|------------|------|-----|-------|------------|

0 rows × 21 columns

Check Assumptions

1. Linearity - Needs a linear relationship between x and y to get linear regression
2. Normality - Residuals are normally distributed
 - Check after modeling
3. No multicollinearity = independant variables not highly correlated
4. Homoscedasticity - variance of error terms are similar across the values of the independant variables
 - Check after modeling


```
In [98]: #Let's create scatterplots of all our data to get an idea linear relationships
#We can also use this to check our outliers
def scatter_view(data_set):
    ncols = 3
    nrows = int(np.ceil(len(data_set.columns) / (1.0*ncols)))
    fig, axes = plt.subplots(nrows=nrows, ncols=ncols, sharey=True, figsize=
    axe = axes.ravel()
    for idx, column in enumerate(data_set.columns):
        data_set.plot(kind='scatter', x=column, y='price', ax=axe[idx], lab
```

```
In [99]: scatter_view(data)
```




Some takeaways

- Our outliers look mostly mitigated, especially in the columns we are most concerned about
- id, date, yr_built, yr_renovated, zipcode, lat, long and sqft_lot15 don't seem to have much of a correlation

```
In [100]: #Heatmap to double check
plt.figure(figsize=(20,20))
matrix = np.triu(data.corr())
sns.heatmap(data.corr(), annot=True, mask=matrix)
```

Out[100]: <AxesSubplot:>



With this heatmap we are simply looking for any bit of correlation along the price

- 1 is perfect correltaion
- .5 is medium

- .3 and less is small to nil
- 0 is no correlation

While this would not be exactly normal we are going to set an alpha level of .2 for our correlation, anything less will be tossed out, anything more will be kept. Lost features - id, date (is an object and won't show here), sqft_lot, condition, yr_built, yr_renovated, long, lat (while it does correlate, it doesn't really do us much good if we don't have long), sqft_lot15

We will additionally remove view, as our business problem is trying to help people find a price to list their homes at there is no way for them to know how many views they will get.

```
In [101]: unnneeded columns to make data easier to digest, with reasoning
drop(['id', 'date', 'view', 'sqft_lot', 'condition', 'yr_built', 'yr_renovated',
      'long', 'lat', 'sqft_lot15'])
```

Out[101]:

| | price | bedrooms | bathrooms | sqft_living | floors | waterfront | grade | sqft_above | sqft_basement |
|---|----------|----------|-----------|-------------|--------|------------|-------|------------|---------------|
| 0 | 221900.0 | 3 | 1.0 | 1180 | 1.0 | 0.0 | 7 | 1180 | 0.0 |
| 1 | 538000.0 | 3 | 2.0 | 2570 | 2.0 | 0.0 | 7 | 2170 | 400.0 |
| 2 | 180000.0 | 2 | 1.0 | 770 | 1.0 | 0.0 | 6 | 770 | 0.0 |
| 3 | 604000.0 | 4 | 3.0 | 1960 | 1.0 | 0.0 | 7 | 1050 | 910.0 |
| 4 | 510000.0 | 3 | 2.0 | 1680 | 1.0 | 0.0 | 8 | 1680 | 0.0 |

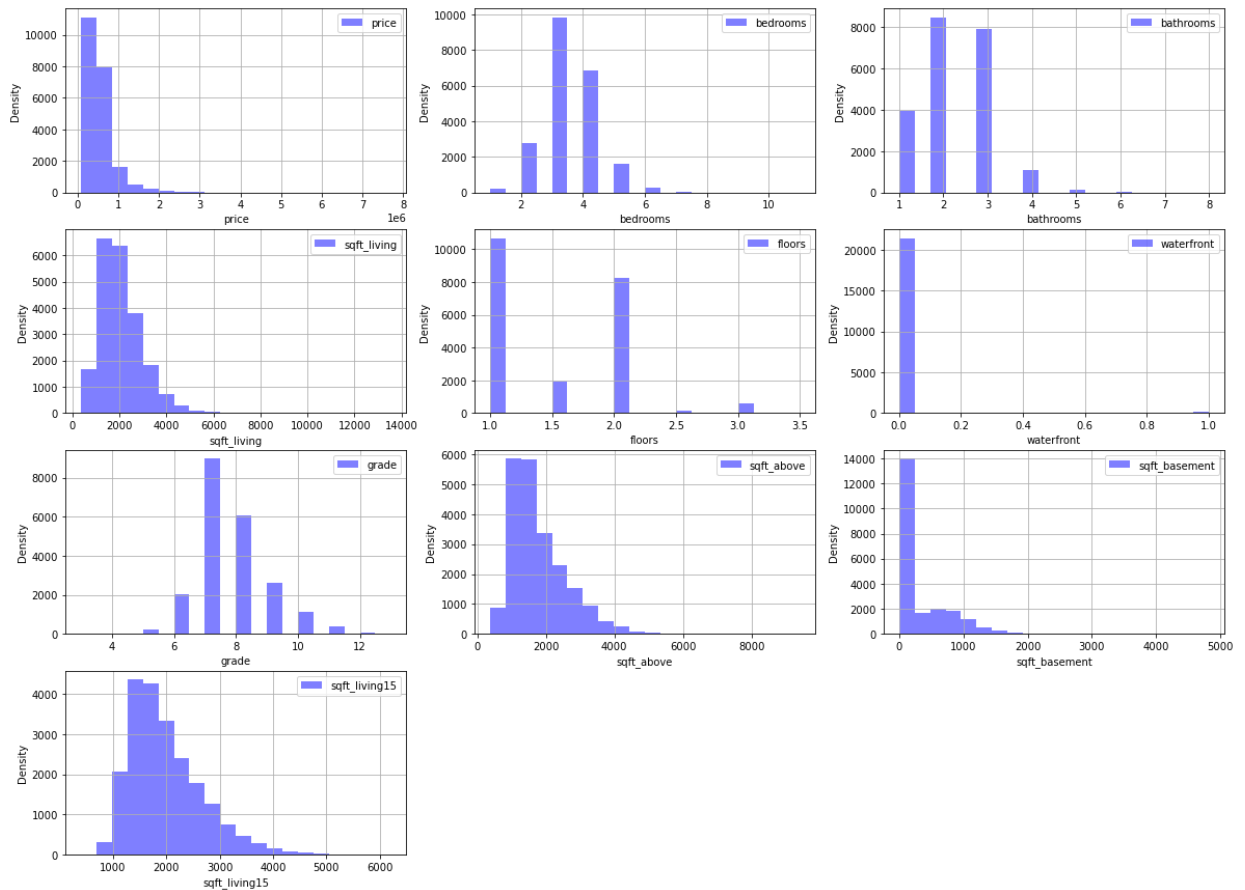
Reasoning for removing columns

- id
 - Now that we have checked duplicates we no longer really need the id as there is no correlation
- date
 - Will be difficult to put into a linear model, as well as be hard to have any use going into the future
- view
 - Does not fit our business problem
- sqft_lot and sqft_lot15
 - Looking at correlations it does not appear that lot has much of an impact on selling price
 - Mind you I am sure this is all location based, and if we had more time to refine the model we could reinclude this.
- yr_built
 - Scatter and heatmap do not show much of a correlation here
- yr_renovated
 - There is a lot of NaN values here, and it will be difficult to select a proper metric for that NaN data that will not end up skewing our model
- zipcode
 - There are greater than 70 zipcodes in the dataset, and for the current project it will be too time consuming to put them into different bins and model on that. Again, could be good

for future modeling

- Likely helpful in further investigations
- lat, long
 - Each home has it's own value, and otherwise similar reasoning as the zipcode column
 - A point to look at for further investigation

```
In [102]: #Lets take an overview of our remaining data
histogram_view(data)
```



Most of the data at least looks somewhat normal, at least we should be able to standardize and such to get better results

Dealing with the massive tail of sqft_basement

```
In [103]: #sqft_basement has many values that are 0, going to create a new column tha
#for no basement, and 1 for has basement (which will be determined as 0 sqf
#a basement)
conditions = [(data['sqft_basement']==0), (data['sqft_basement'] > 0)]
values = [0,1]
data['basement'] = np.select(conditions, values)
data.head()
```

Out[103]:


| | price | bedrooms | bathrooms | sqft_living | floors | waterfront | grade | sqft_above | sqft_basemen |
|---|----------|----------|-----------|-------------|--------|------------|-------|------------|--------------|
| 0 | 221900.0 | 3 | 1.0 | 1180 | 1.0 | 0.0 | 7 | 1180 | 0.0 |
| 1 | 538000.0 | 3 | 2.0 | 2570 | 2.0 | 0.0 | 7 | 2170 | 400.0 |
| 2 | 180000.0 | 2 | 1.0 | 770 | 1.0 | 0.0 | 6 | 770 | 0.0 |
| 3 | 604000.0 | 4 | 3.0 | 1960 | 1.0 | 0.0 | 7 | 1050 | 910.0 |
| 4 | 510000.0 | 3 | 2.0 | 1680 | 1.0 | 0.0 | 8 | 1680 | 0.0 |

```
In [104]: #Drop obsolete sqft_basement column
data.drop(['sqft_basement'], axis=1, inplace=True)
#Sanity check
data.columns
```

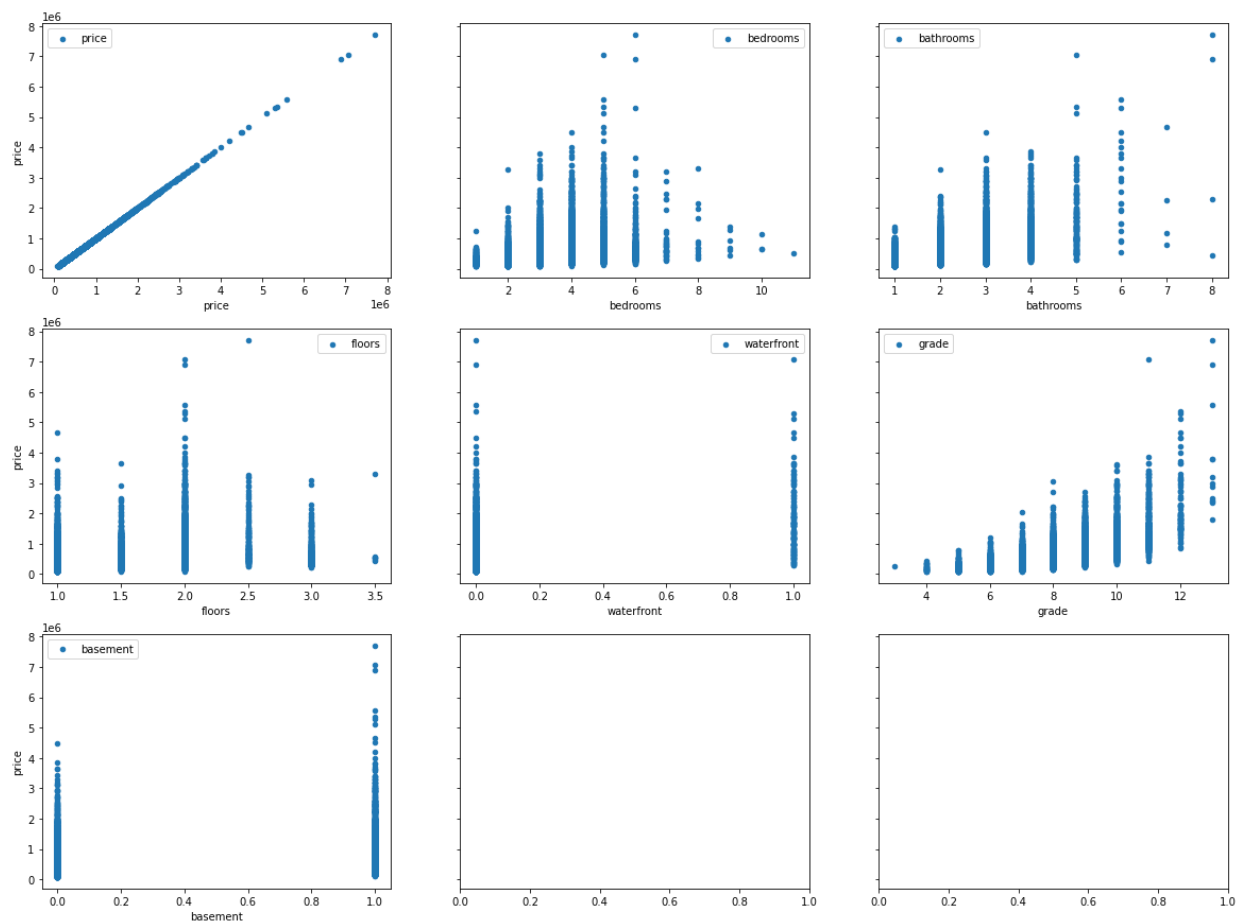
Out[104]: Index(['price', 'bedrooms', 'bathrooms', 'sqft_living', 'floors', 'waterf
ront',
 'grade', 'sqft_above', 'sqft_living15', 'basement'],
 dtype='object')

Dealing with categorical data

```
In [105]: #Check out histograms and Scatter plots  
categoricals = ['bedrooms', 'bathrooms', 'floors', 'waterfront', 'grade', 'basement']  
histogram_view(data[categoricals])
```



```
In [106]: #Add price back in to cat_data to
categoricals_w_price = ['price', 'bedrooms', 'bathrooms', 'floors', 'waterfr
scatter_view(data[categoricals_w_price])
```



Looks like most of our conditionals have some sort of solid relationship with our price.

##ONE HOT ENCODING

```
In [107]: #Check data types to get ready to O-H-E
#We had an issue with these not being in category when running this before
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21593 entries, 0 to 21596
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   price                 21593 non-null  float64
1   bedrooms              21593 non-null  int64
2   bathrooms              21593 non-null  float64
3   sqft_living           21593 non-null  int64
4   floors                21593 non-null  float64
5   waterfront            21593 non-null  float64
6   grade                 21593 non-null  int64
7   sqft_above            21593 non-null  int64
8   sqft_living15         21593 non-null  int64
9   basement              21593 non-null  int64
dtypes: float64(4), int64(6)
memory usage: 1.8 MB
```

```
In [108]: cat_data = data[categoricals].astype('category')
```

```
In [109]: cat_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21593 entries, 0 to 21596
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   bedrooms              21593 non-null  category
1   bathrooms              21593 non-null  category
2   floors                21593 non-null  category
3   waterfront            21593 non-null  category
4   grade                 21593 non-null  category
5   basement              21593 non-null  category
dtypes: category(6)
memory usage: 296.8 KB
```


In [110]: *#Looks like bedrooms, bathrooms, floors, waterfront, condition, and basements*
#Let's turn them into dummies for the future modeling

```
dummies = pd.get_dummies(cat_data[categoricals], prefix=categoricals, drop_
data_preprocessed = data.drop(categoricals, axis=1)
data_preprocessed = pd.concat([data_preprocessed, dummies], axis=1)
data_preprocessed.head()
```

Out[110]:

| | price | sqft_living | sqft_above | sqft_living15 | bedrooms_2 | bedrooms_3 | bedrooms_4 | bedrooms |
|---|----------|-------------|------------|---------------|------------|------------|------------|----------|
| 0 | 221900.0 | 1180 | 1180 | 1340 | 0 | 1 | 0 | |
| 1 | 538000.0 | 2570 | 2170 | 1690 | 0 | 1 | 0 | |
| 2 | 180000.0 | 770 | 770 | 2720 | 1 | 0 | 0 | |
| 3 | 604000.0 | 1960 | 1050 | 1360 | 0 | 0 | 1 | |
| 4 | 510000.0 | 1680 | 1680 | 1800 | 0 | 1 | 0 | |

5 rows × 38 columns

```
In [111]: #One final look at our preprocessed data
data_preprocessed.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21593 entries, 0 to 21596
Data columns (total 38 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 21593 non-null  float64
1   sqft_living           21593 non-null  int64
2   sqft_above            21593 non-null  int64
3   sqft_living15         21593 non-null  int64
4   bedrooms_2           21593 non-null  uint8
5   bedrooms_3           21593 non-null  uint8
6   bedrooms_4           21593 non-null  uint8
7   bedrooms_5           21593 non-null  uint8
8   bedrooms_6           21593 non-null  uint8
9   bedrooms_7           21593 non-null  uint8
10  bedrooms_8           21593 non-null  uint8
11  bedrooms_9           21593 non-null  uint8
12  bedrooms_10          21593 non-null  uint8
13  bedrooms_11          21593 non-null  uint8
14  bathrooms_2.0        21593 non-null  uint8
15  bathrooms_3.0        21593 non-null  uint8
16  bathrooms_4.0        21593 non-null  uint8
17  bathrooms_5.0        21593 non-null  uint8
18  bathrooms_6.0        21593 non-null  uint8
19  bathrooms_7.0        21593 non-null  uint8
20  bathrooms_8.0        21593 non-null  uint8
21  floors_1.5           21593 non-null  uint8
22  floors_2.0           21593 non-null  uint8
23  floors_2.5           21593 non-null  uint8
24  floors_3.0           21593 non-null  uint8
25  floors_3.5           21593 non-null  uint8
26  waterfront_1.0       21593 non-null  uint8
27  grade_4              21593 non-null  uint8
28  grade_5              21593 non-null  uint8
29  grade_6              21593 non-null  uint8
30  grade_7              21593 non-null  uint8
31  grade_8              21593 non-null  uint8
32  grade_9              21593 non-null  uint8
33  grade_10             21593 non-null  uint8
34  grade_11             21593 non-null  uint8
35  grade_12             21593 non-null  uint8
36  grade_13             21593 non-null  uint8
37  basement_1           21593 non-null  uint8
dtypes: float64(1), int64(3), uint8(34)
memory usage: 1.5 MB
```

Begin modeling

We're going to start modeling with just one simple model that has had no scaling or transformations done on it. We're also going to make a function that returns all

```

In [112]: def linreg_process(X_data, y_data):
    """ This is a function that returns all relevant data from a modeling p
    All you need is to know what you x and y variables in a dataset are and
    them into test and train sets, create an OLS model and print the summar
    tests and print out those values"""
    linreg = LinearRegression()
    X_train, X_test, y_train, y_test = train_test_split(X_data, y_data, tes
    predictors = sm.add_constant(X_train)
    model = sm.OLS(y_train, predictors).fit()
    print(model.summary())
    linreg.fit(X_train, y_train)
    y_pred_train = linreg.predict(X_train)
    y_pred_test = linreg.predict(X_test)
    print('----- Scoring Stats -----')
    print(f"R2 Train Score: {r2_score(y_train, y_pred_train)}")
    print(f"R2 Test Score: {r2_score(y_test, y_pred_test)}")
    print(f"Intercept: {linreg.intercept_}")
    print(f"Coefficient: {linreg.coef_}")
    print(f"MSE Train: {mean_squared_error(y_train, y_pred_train)}")
    print(f"MSE Test: {mean_squared_error(y_test, y_pred_test)}")
    residuals = model.resid
    print('----- Q-Q Plot -----')
    fig = sm.graphics.qqplot(residuals, dist=stats.norm, line='45',fit=True)
    fig.show()

```

```
In [113]: X = data_preprocessed.drop('price', axis=1)
y = data_preprocessed['price']
linreg_process(X, y)
```

OLS Regression Results

```
=====
=====
Dep. Variable:          price    R-squared:
0.647
Model:                  OLS      Adj. R-squared:
0.646
Method:                 Least Squares    F-statistic:
823.5
Date:                   Wed, 03 Feb 2021    Prob (F-statistic):
0.00
Time:                   20:27:22    Log-Likelihood:          -2.220
3e+05
No. Observations:      16194    AIC:          4.44
1e+05
Df Residuals:          16157    BIC:          4.44
4e+05
Df Model:              36
Covariance Type:       nonrobust
=====
=====
```

| | coef | std err | t | P> t | [0.025 |
|---------------|------------|----------|--------|-------|-----------|
| 0.975] | | | | | |
| ----- | | | | | |
| const | 5.027e+05 | 2.06e+04 | 24.406 | 0.000 | 4.62e+05 |
| 5.43e+05 | | | | | |
| sqft_living | 164.1925 | 7.616 | 21.558 | 0.000 | 149.264 |
| 179.121 | | | | | |
| sqft_above | -38.4451 | 8.421 | -4.565 | 0.000 | -54.952 |
| -21.938 | | | | | |
| sqft_living15 | 45.8615 | 4.259 | 10.768 | 0.000 | 37.513 |
| 54.210 | | | | | |
| bedrooms_2 | 7681.1604 | 1.86e+04 | 0.414 | 0.679 | -2.87e+04 |
| 4.41e+04 | | | | | |
| bedrooms_3 | -4.74e+04 | 1.85e+04 | -2.557 | 0.011 | -8.37e+04 |
| -1.11e+04 | | | | | |
| bedrooms_4 | -6.733e+04 | 1.89e+04 | -3.554 | 0.000 | -1.04e+05 |
| -3.02e+04 | | | | | |
| bedrooms_5 | -5.366e+04 | 2e+04 | -2.677 | 0.007 | -9.3e+04 |
| -1.44e+04 | | | | | |
| bedrooms_6 | -6.44e+04 | 2.43e+04 | -2.648 | 0.008 | -1.12e+05 |
| -1.67e+04 | | | | | |
| bedrooms_7 | -1.495e+05 | 4.41e+04 | -3.391 | 0.001 | -2.36e+05 |
| -6.31e+04 | | | | | |
| bedrooms_8 | 1.411e+05 | 7.26e+04 | 1.944 | 0.052 | -1201.447 |
| 2.83e+05 | | | | | |
| bedrooms_9 | -4.852e+05 | 1.15e+05 | -4.211 | 0.000 | -7.11e+05 |
| -2.59e+05 | | | | | |
| bedrooms_10 | -8.833e+04 | 1.56e+05 | -0.566 | 0.571 | -3.94e+05 |
| 2.17e+05 | | | | | |
| bedrooms_11 | -5.731e+04 | 2.19e+05 | -0.262 | 0.794 | -4.87e+05 |
| 3.72e+05 | | | | | |

| | | | | | |
|----------------|------------|----------|-------------------|-------|-----------|
| bathrooms_2.0 | -3.078e+04 | 5947.522 | -5.175 | 0.000 | -4.24e+04 |
| -1.91e+04 | | | | | |
| bathrooms_3.0 | -7.117e+04 | 7542.674 | -9.436 | 0.000 | -8.6e+04 |
| -5.64e+04 | | | | | |
| bathrooms_4.0 | -8932.8997 | 1.19e+04 | -0.752 | 0.452 | -3.22e+04 |
| 1.43e+04 | | | | | |
| bathrooms_5.0 | 1.898e+05 | 2.43e+04 | 7.811 | 0.000 | 1.42e+05 |
| 2.37e+05 | | | | | |
| bathrooms_6.0 | 7.008e+05 | 6.03e+04 | 11.612 | 0.000 | 5.82e+05 |
| 8.19e+05 | | | | | |
| bathrooms_7.0 | -9.916e+04 | 1.29e+05 | -0.767 | 0.443 | -3.52e+05 |
| 1.54e+05 | | | | | |
| bathrooms_8.0 | 1.544e+06 | 1.25e+05 | 12.391 | 0.000 | 1.3e+06 |
| 1.79e+06 | | | | | |
| floors_1.5 | 1.083e+05 | 6362.073 | 17.023 | 0.000 | 9.58e+04 |
| 1.21e+05 | | | | | |
| floors_2.0 | -7263.3548 | 5444.174 | -1.334 | 0.182 | -1.79e+04 |
| 3407.830 | | | | | |
| floors_2.5 | 1.898e+05 | 2.03e+04 | 9.348 | 0.000 | 1.5e+05 |
| 2.3e+05 | | | | | |
| floors_3.0 | 7.993e+04 | 1.13e+04 | 7.054 | 0.000 | 5.77e+04 |
| 1.02e+05 | | | | | |
| floors_3.5 | 2.173e+05 | 8.34e+04 | 2.606 | 0.009 | 5.38e+04 |
| 3.81e+05 | | | | | |
| waterfront_1.0 | 7.259e+05 | 2.19e+04 | 33.177 | 0.000 | 6.83e+05 |
| 7.69e+05 | | | | | |
| grade_4 | -4.297e+05 | 4.48e+04 | -9.587 | 0.000 | -5.18e+05 |
| -3.42e+05 | | | | | |
| grade_5 | -4.571e+05 | 1.81e+04 | -25.298 | 0.000 | -4.93e+05 |
| -4.22e+05 | | | | | |
| grade_6 | -4.152e+05 | 1.21e+04 | -34.417 | 0.000 | -4.39e+05 |
| -3.92e+05 | | | | | |
| grade_7 | -3.564e+05 | 1.08e+04 | -32.899 | 0.000 | -3.78e+05 |
| -3.35e+05 | | | | | |
| grade_8 | -2.763e+05 | 1.06e+04 | -25.967 | 0.000 | -2.97e+05 |
| -2.55e+05 | | | | | |
| grade_9 | -1.354e+05 | 1.1e+04 | -12.274 | 0.000 | -1.57e+05 |
| -1.14e+05 | | | | | |
| grade_10 | 3.657e+04 | 1.23e+04 | 2.964 | 0.003 | 1.24e+04 |
| 6.08e+04 | | | | | |
| grade_11 | 2.83e+05 | 1.66e+04 | 17.025 | 0.000 | 2.5e+05 |
| 3.16e+05 | | | | | |
| grade_12 | 6.223e+05 | 2.78e+04 | 22.401 | 0.000 | 5.68e+05 |
| 6.77e+05 | | | | | |
| grade_13 | 1.631e+06 | 7.48e+04 | 21.800 | 0.000 | 1.48e+06 |
| 1.78e+06 | | | | | |
| basement_1 | 4.404e+04 | 6345.616 | 6.940 | 0.000 | 3.16e+04 |
| 5.65e+04 | | | | | |
| ===== | | | | | |
| ===== | | | | | |
| Omnibus: | | 7897.630 | Durbin-Watson: | | |
| 1.992 | | | | | |
| Prob(Omnibus): | | 0.000 | Jarque-Bera (JB): | | 20886 |
| 2.528 | | | | | |
| Skew: | | 1.789 | Prob(JB): | | |
| 0.00 | | | | | |
| Kurtosis: | | 20.226 | Cond. No. | | 6.5 |

9e+15

=====

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 4.9e-21. This might indicate that there are

strong multicollinearity problems or that the design matrix is singular.

----- Scoring Stats -----

R2 Train Score: 0.6472383143454867

R2 Test Score: 0.6483069411223183

Intercept: 552948.0719841787

Coefficient: [1.64192500e+02 -3.84451339e+01 4.58614630e+01 7.68116040e+03

-4.74003750e+04 -6.73257548e+04 -5.36625632e+04 -6.44042498e+04

-1.49450619e+05 1.41085739e+05 -4.85187934e+05 -8.83330269e+04

-5.73074216e+04 -3.07789575e+04 -7.11689597e+04 -8.93289970e+03

1.89752547e+05 7.00755831e+05 -9.91557343e+04 1.54420159e+06

1.08303477e+05 -7.26335475e+03 1.89827523e+05 7.99276845e+04

2.17292837e+05 7.25886994e+05 -4.80007439e+05 -5.07366034e+05

-4.65455409e+05 -4.06682187e+05 -3.26603941e+05 -1.85626860e+05

-1.36942190e+04 2.32692519e+05 5.72009269e+05 1.58073430e+06

4.40413935e+04]


MSE Train: 47490846003.166756

MSE Test: 47737041399.685776

----- Q-Q Plot -----

<ipython-input-112-dbf471e5a3b>:24: UserWarning: Matplotlib is currently using module://ipykernel.pylab.backend_inline, which is a non-GUI backend, so cannot show the figure.

fig.show()



Some takeaways from the model

- R2 Score:

- The training got about a .647 meaning that the model explains roughly 65% of the variability of the response data around the mean
- This is not a great score, but certainly something we can work with
- When comparing the train and test R2 scores it appears that the test R2 is slightly higher at .648 this seems to indicate some very mild overfitting, but this could also just be randomness in the samples (cross validation should show if this is the case or not)
- P-values:
 - Most of our data is below of the alpha level 0.05, there are quite a few outliers (mostly within some of our one-hot-encoded variables that are certainly skewing the data somewhat
- Mean Squared Error (MSE):
 - This is an extremely high score, this indicates that we have a lot of errors in our modeling and this thing is absolutely no where close to perfect
 - Scaling and transformations should help
- JB Test
 - The JB test indicates that our data is anything but normal (backed up by the QQ plot)
- F stat
 - Indicates that our data is not homoscedastic
- Overall, our data needs a lot of work to get anywhere close to a respectable model

Some takeaways from the Q-Q Plot

- We do not have a normal distribution
- looks like a whole lot of our data is heavy on the left, but performing some scaling should help a lot with fixing this.

Quick Cross-Validation

```
In [114]: #This will give us the mean Mean Squared Error to check our model
def cross_val(X_data, y_data):
    linreg = LinearRegression()
    X_train_cv, X_test_cv, y_train_cv, y_test_cv = train_test_split(X_data,
        predictors = sm.add_constant(X_train_cv)
    linreg.fit(X_train_cv, y_train_cv)
    mse = make_scorer(mean_squared_error)
    cv_5_results = cross_val_score(linreg, X_data, y_data, cv=5, scoring=mse)
    print(cv_5_results.mean())
```

```
In [115]: cross_val(X, y)

48807152240.95346
```

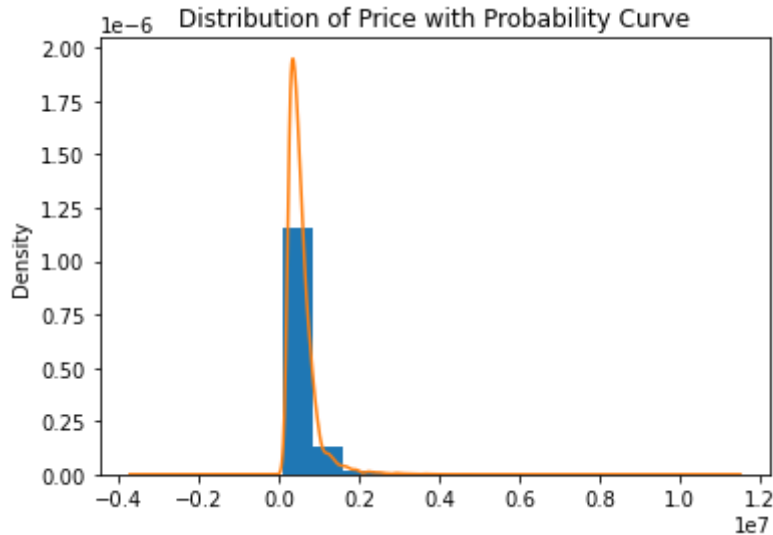
Not all that far off from our data, we have a lot of work to get this running better

Scaling

Let's start the scaling process by taking a look at our most important variable, the target variable price!

```
In [116]: #Plot histogram with probability density to check distribution of y axis
plt.hist(data_preprocessed['price'], density = True)
data_preprocessed['price'].plot.kde()
plt.title("Distribution of Price with Probability Curve")
```

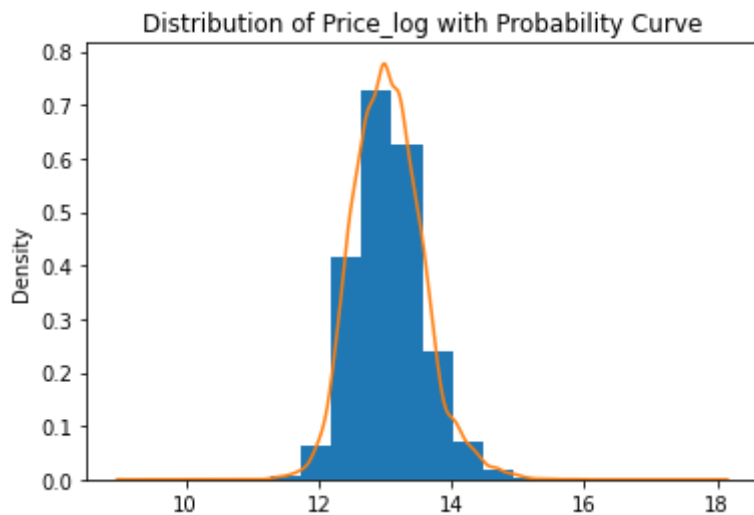
```
Out[116]: Text(0.5, 1.0, 'Distribution of Price with Probability Curve')
```



This shows that our data is heavily skewed, and one of the easiest ways to "normalize" this is to perform a log-transformation.

```
In [117]: #Now lets start our scaling with normalizing this with a log-transform
data_preprocessed['price_log'] = np.log(data['price'])
plt.hist(data_preprocessed['price_log'], density = True)
data_preprocessed['price_log'].plot.kde()
plt.title("Distribution of Price_log with Probability Curve")
```

```
Out[117]: Text(0.5, 1.0, 'Distribution of Price_log with Probability Curve')
```



Woo! look at that normal data. now let's get to the independant variables

Standard Scaler


We're going to start with using a standard scaler on our data. This can be sensitive to outliers, but as we have already mitigated those I am not too concerned about that, what it will do is bring all of our data to a mean of 0 with a standard deviation of 1. This centers the data and brings it to a Gaussian standard normal deviation, which helps with any modeling.

```
In [118]: #Standard scaler
```

```
scaler = StandardScaler()
dataX2 = data_preprocessed.drop(['price_log', 'price'], axis=1)
y2 = data_preprocessed['price_log']
X_scaled = scaler.fit_transform(dataX2)
scaled_df = pd.DataFrame(X_scaled, columns = dataX2.columns) #For later use
```

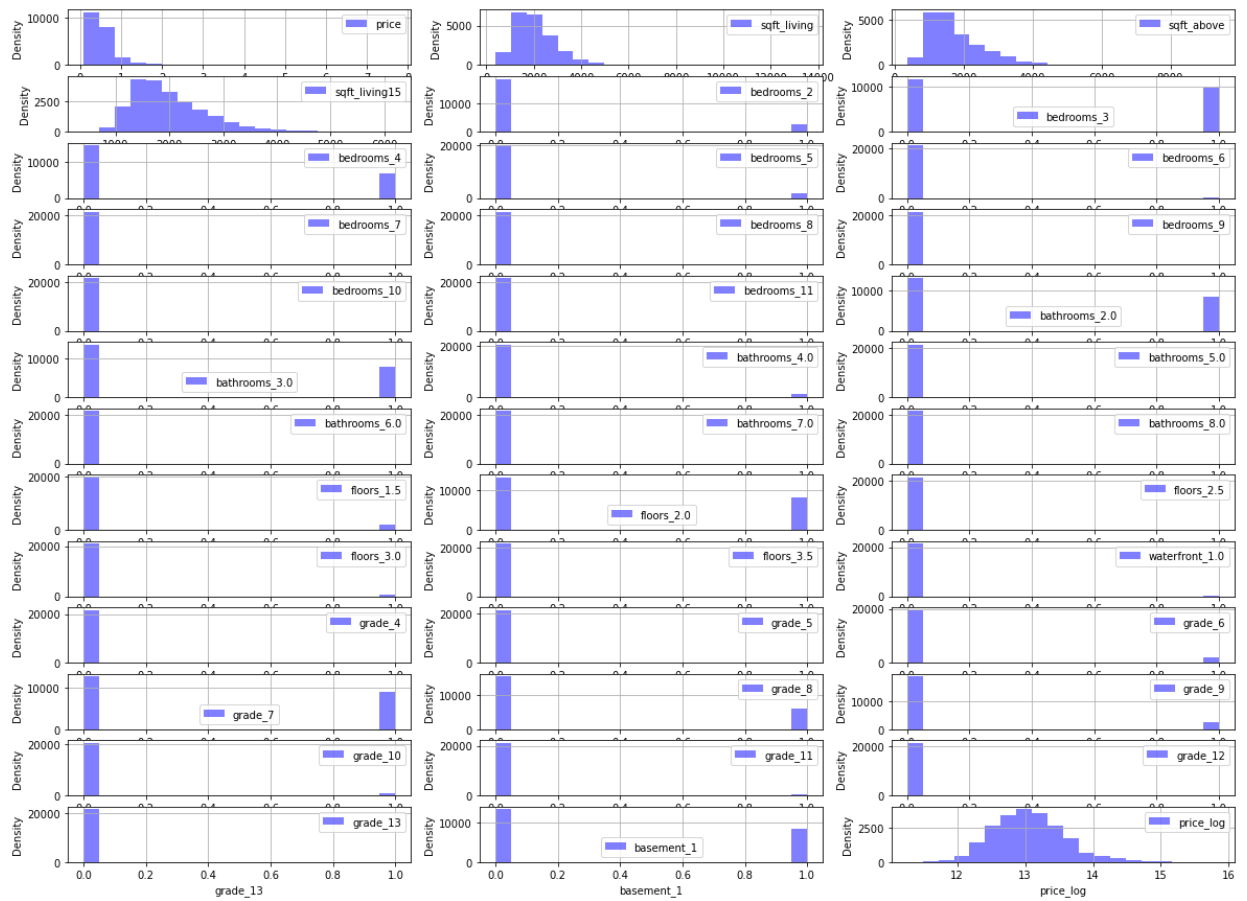
```
In [119]: cv_scaled_df = pd.concat([scaled_df, data_preprocessed['price_log']], axis=1)
```

```
In [120]: histogram_view(scaled_df)
```



Compared with the original

```
In [121]: histogram_view(data_preprocessed)
```



All in all the data is now centered around 0 and looks to be slightly more normal, hopefully we can see the results for real on our model.

Model #2 log-transformed y and scaled X

```
In [122]: linreg_process(X_scaled, y2)
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price_log    R-squared:
0.605
Model:                  OLS    Adj. R-squared:
0.605
Method:                Least Squares    F-statistic:
688.7
Date:                  Wed, 03 Feb 2021    Prob (F-statistic):
0.00
Time:                  20:27:31    Log-Likelihood:          -5
032.3
No. Observations:      16194    AIC:                  1.01
4e+04
Df Residuals:          16157    BIC:                  1.04
2e+04
Df Model:              36
Covariance Type:      nonrobust
=====
=====
                        coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const          13.0458      0.003    5020.401      0.000      13.041      1
3.051
x1              0.1605      0.011     15.161      0.000       0.140
0.181
x2            -0.0060      0.011     -0.573      0.567      -0.027
0.015
x3              0.0787      0.004     17.790      0.000       0.070
0.087
x4              0.0112      0.009       1.198      0.231      -0.007
0.030
x5            -0.0401      0.014     -2.868      0.004      -0.068      -
0.013
x6            -0.0428      0.013     -3.201      0.001      -0.069      -
0.017
x7            -0.0203      0.008     -2.548      0.011      -0.036      -
0.005
x8            -0.0086      0.004     -2.092      0.036      -0.017      -
0.001
x9            -0.0030      0.003     -1.073      0.283      -0.008
0.002
x10             0.0022      0.003       0.811      0.417      -0.003
0.007
x11             0.0025      0.003       0.868      0.385      -0.003
0.008
x12             0.0011      0.003       0.389      0.697      -0.004
0.007
x13             0.0004      0.002       0.166      0.868      -0.004
0.005
x14            -0.0016      0.004     -0.353      0.724      -0.010
0.007

```

| | | | | | | |
|----------------|---------|-------|-------------------|-------|--------|-----|
| x15 | -0.0238 | 0.006 | -4.332 | 0.000 | -0.035 | - |
| 0.013 | | | | | | |
| x16 | 0.0025 | 0.004 | 0.643 | 0.520 | -0.005 | |
| 0.010 | | | | | | |
| x17 | 0.0044 | 0.003 | 1.397 | 0.162 | -0.002 | |
| 0.010 | | | | | | |
| x18 | 0.0024 | 0.003 | 0.852 | 0.394 | -0.003 | |
| 0.008 | | | | | | |
| x19 | -0.0083 | 0.003 | -3.126 | 0.002 | -0.014 | - |
| 0.003 | | | | | | |
| x20 | -0.0060 | 0.003 | -2.343 | 0.019 | -0.011 | - |
| 0.001 | | | | | | |
| x21 | 0.0654 | 0.003 | 23.880 | 0.000 | 0.060 | |
| 0.071 | | | | | | |
| x22 | 0.0018 | 0.004 | 0.461 | 0.645 | -0.006 | |
| 0.010 | | | | | | |
| x23 | 0.0162 | 0.003 | 6.133 | 0.000 | 0.011 | |
| 0.021 | | | | | | |
| x24 | 0.0281 | 0.003 | 9.855 | 0.000 | 0.022 | |
| 0.034 | | | | | | |
| x25 | 0.0040 | 0.002 | 1.743 | 0.081 | -0.000 | |
| 0.008 | | | | | | |
| x26 | 0.0456 | 0.003 | 16.795 | 0.000 | 0.040 | |
| 0.051 | | | | | | |
| x27 | -0.0196 | 0.003 | -7.552 | 0.000 | -0.025 | - |
| 0.015 | | | | | | |
| x28 | -0.0570 | 0.003 | -21.639 | 0.000 | -0.062 | - |
| 0.052 | | | | | | |
| x29 | -0.0994 | 0.003 | -34.106 | 0.000 | -0.105 | - |
| 0.094 | | | | | | |
| x30 | -0.0682 | 0.002 | -30.110 | 0.000 | -0.073 | - |
| 0.064 | | | | | | |
| x31 | 0.0230 | 0.002 | 11.322 | 0.000 | 0.019 | |
| 0.027 | | | | | | |
| x32 | 0.0849 | 0.003 | 31.943 | 0.000 | 0.080 | |
| 0.090 | | | | | | |
| x33 | 0.0901 | 0.003 | 30.225 | 0.000 | 0.084 | |
| 0.096 | | | | | | |
| x34 | 0.0667 | 0.003 | 21.001 | 0.000 | 0.060 | |
| 0.073 | | | | | | |
| x35 | 0.0373 | 0.003 | 12.865 | 0.000 | 0.032 | |
| 0.043 | | | | | | |
| x36 | 0.0191 | 0.003 | 6.242 | 0.000 | 0.013 | |
| 0.025 | | | | | | |
| x37 | 0.0643 | 0.005 | 13.739 | 0.000 | 0.055 | |
| 0.073 | | | | | | |
| ===== | | | | | | |
| ===== | | | | | | |
| Omnibus: | | 0.714 | Durbin-Watson: | | | |
| 1.998 | | | | | | |
| Prob(Omnibus): | | 0.700 | Jarque-Bera (JB): | | | |
| 0.687 | | | | | | |
| Skew: | | 0.011 | Prob(JB): | | | |
| 0.709 | | | | | | |
| Kurtosis: | | 3.024 | Cond. No. | | | 2.4 |
| 9e+15 | | | | | | |
| ===== | | | | | | |

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 1.26e-26. This might indicate that there are

strong multicollinearity problems or that the design matrix is singular.

----- Scoring Stats -----

R2 Train Score: 0.6054593246510145

R2 Test Score: -6.224978467286221e+17

Intercept: -1420043.7427191257

Coefficient: [1.60479524e-01 -6.04218143e-03 7.86501653e-02 1.12362217e-02

-4.01157409e-02 -4.28253718e-02 -2.02757531e-02 -8.59895068e-03

-2.99946278e-03 2.18514193e-03 2.51684804e-03 1.08169522e-03

3.76149136e-04 -1.55096664e-03 -2.38491002e-02 2.54393797e-03

4.36799491e-03 2.42034242e-03 -8.33044592e-03 -6.02156888e-03

6.53613765e-02 1.84311712e-03 1.62306068e-02 2.80546470e-02

3.96455921e-03 4.56223927e-02 1.08360848e+09 3.22791809e+09

8.96274914e+09 1.51113834e+10 1.37809411e+10 1.00038410e+10

6.83997384e+09 4.11934268e+09 1.96453758e+09 7.52147660e+08

6.42681412e-02]


MSE Train: 0.10900365141648688

MSE Test: 1.741502357089267e+17

----- Q-Q Plot -----

<ipython-input-112-dbf471e5a3b>:24: UserWarning: Matplotlib is currently using module://ipykernel.pylab.backend_inline, which is a non-GUI backend, so cannot show the figure.

fig.show()



- R2 Score:
 - Pretty much the same as our first model, nothing much to report here
- P-values:
 - Same issue as before, some of our one_hot_encoded variables are certainly skewing the data some here
- Mean Squared Error (MSE):
 - This is where we saw the best effect on our scaling, as our data is much more normal now, there is not a lot of error in our model, it's just not that good at predicting
- JB Test
 - Massive change here, we have a respectable score showing that our data does not have a lot of skew or kurtosis, it is much closer to normal than before the scaling
- F stat
 - Indicates that our data is not homoscedastic, we have errors all over the regression line, but it is better than the previous attempt
- Overall, I think we moved in the right direction, but there still isn't much predictive power with this model

Some takeaways from the Q-Q Plot

- With our scaling it looks like we have achieved the goal of a normal distribution of our residuals!

```
In [123]: #Quick cross val for reference
cross_val(X_scaled, y2)
```

5.073375499028728e+16

This is actually considerably higher than in the tests we ran, we have more variance from the mean than initially expected.

Feature Selection

This should help the model by showing us which features from the independent data have the highest

```
In [124]: #We are going to remove the worst 7 variables from our preprocessed data, 1
#impactful independent variables

lr = LinearRegression()
selector = RFE(lr, n_features_to_select = 30)
selector = selector.fit(X, y.ravel()) # convert y to 1d np array to prevent
selector.support_
selected_columns = X.columns[selector.support_ ]
```

```
In [125]: selected_columns
```

```
Out[125]: Index(['bedrooms_2', 'bedrooms_4', 'bedrooms_5', 'bedrooms_6', 'bedrooms_7',  
                'bedrooms_8', 'bedrooms_9', 'bedrooms_10', 'bedrooms_11',  
                'bathrooms_4.0', 'bathrooms_5.0', 'bathrooms_6.0', 'bathrooms_7.0',  
                'bathrooms_8.0', 'floors_1.5', 'floors_2.5', 'floors_3.0', 'floors_3.5',  
                'waterfront_1.0', 'grade_4', 'grade_5', 'grade_6', 'grade_7', 'grade_8',  
                'grade_9', 'grade_10', 'grade_11', 'grade_12', 'grade_13',  
                'basement_1'],  
               dtype='object')
```

```
In [126]: linreg_process(X[selected_columns], y)
```

```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:
0.613
Model:                          OLS    Adj. R-squared:
0.612
Method:                        Least Squares    F-statistic:
882.8
Date:                          Wed, 03 Feb 2021    Prob (F-statistic):
0.00
Time:                          20:27:32    Log-Likelihood:            -2.227
8e+05
No. Observations:              16194    AIC:                        4.45
6e+05
Df Residuals:                  16164    BIC:                        4.45
9e+05
Df Model:                      29
Covariance Type:               nonrobust
=====
=====
                                coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const                8.158e+05    1.03e+04     79.439    0.000    7.96e+05
8.36e+05
bedrooms_2           1.915e+04    5900.524     3.245    0.001    7583.276
3.07e+04
bedrooms_4           2.919e+04    4329.057     6.743    0.000    2.07e+04
3.77e+04
bedrooms_5           8.581e+04    7531.666    11.393    0.000     7.1e+04
1.01e+05
bedrooms_6           9.728e+04    1.61e+04     6.060    0.000    6.58e+04
1.29e+05
bedrooms_7           5.224e+04    4.15e+04     1.259    0.208   -2.91e+04
1.34e+05
bedrooms_8           3.692e+05    7.32e+04     5.044    0.000    2.26e+05
5.13e+05
bedrooms_9           -4.595e+05    1.19e+05    -3.859    0.000   -6.93e+05
-2.26e+05
bedrooms_10          1.134e+05    1.62e+05     0.700    0.484   -2.04e+05
4.31e+05
bedrooms_11          9.41e+04    2.29e+05     0.412    0.680   -3.54e+05
5.42e+05
bathrooms_4.0        1.07e+05    9100.807    11.757    0.000    8.92e+04
1.25e+05
bathrooms_5.0        3.785e+05    2.34e+04    16.156    0.000    3.33e+05
4.24e+05
bathrooms_6.0        9.847e+05    6.19e+04    15.908    0.000    8.63e+05
1.11e+06
bathrooms_7.0        3.369e+05    1.34e+05     2.518    0.012    7.46e+04
5.99e+05
bathrooms_8.0        2.294e+06    1.28e+05    17.903    0.000    2.04e+06
2.55e+06

```


| | | | | | |
|----------------|------------|----------|---------|-------|-----------|
| floors_1.5 | 1.133e+05 | 6335.797 | 17.884 | 0.000 | 1.01e+05 |
| 1.26e+05 | | | | | |
| floors_2.5 | 1.936e+05 | 2.07e+04 | 9.344 | 0.000 | 1.53e+05 |
| 2.34e+05 | | | | | |
| floors_3.0 | 1.431e+04 | 1.1e+04 | 1.299 | 0.194 | -7291.886 |
| 3.59e+04 | | | | | |
| floors_3.5 | 1.798e+05 | 8.72e+04 | 2.063 | 0.039 | 8970.070 |
| 3.51e+05 | | | | | |
| waterfront_1.0 | 7.977e+05 | 2.28e+04 | 34.968 | 0.000 | 7.53e+05 |
| 8.42e+05 | | | | | |
| grade_4 | -6.026e+05 | 4.63e+04 | -13.026 | 0.000 | -6.93e+05 |
| -5.12e+05 | | | | | |
| grade_5 | -6.238e+05 | 1.81e+04 | -34.403 | 0.000 | -6.59e+05 |
| -5.88e+05 | | | | | |
| grade_6 | -5.722e+05 | 1.15e+04 | -49.946 | 0.000 | -5.95e+05 |
| -5.5e+05 | | | | | |
| grade_7 | -4.873e+05 | 1.03e+04 | -47.183 | 0.000 | -5.08e+05 |
| -4.67e+05 | | | | | |
| grade_8 | -3.56e+05 | 1.04e+04 | -34.107 | 0.000 | -3.76e+05 |
| -3.36e+05 | | | | | |
| grade_9 | -1.344e+05 | 1.1e+04 | -12.256 | 0.000 | -1.56e+05 |
| -1.13e+05 | | | | | |
| grade_10 | 1.187e+05 | 1.21e+04 | 9.788 | 0.000 | 9.49e+04 |
| 1.42e+05 | | | | | |
| grade_11 | 4.7e+05 | 1.61e+04 | 29.176 | 0.000 | 4.38e+05 |
| 5.02e+05 | | | | | |
| grade_12 | 9.266e+05 | 2.76e+04 | 33.628 | 0.000 | 8.73e+05 |
| 9.81e+05 | | | | | |
| grade_13 | 2.077e+06 | 7.73e+04 | 26.881 | 0.000 | 1.93e+06 |
| 2.23e+06 | | | | | |
| basement_1 | 9.744e+04 | 3801.342 | 25.634 | 0.000 | 9e+04 |
| 1.05e+05 | | | | | |

=====

=====

| | | | |
|----------------|----------|-------------------|-------|
| Omnibus: | 8884.197 | Durbin-Watson: | |
| 1.986 | | | |
| Prob(Omnibus): | 0.000 | Jarque-Bera (JB): | 27461 |
| 4.996 | | | |
| Skew: | 2.067 | Prob(JB): | |
| 0.00 | | | |
| Kurtosis: | 22.746 | Cond. No. | 9.3 |
| 5e+15 | | | |

=====

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 3e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

----- Scoring Stats -----

R2 Train Score: 0.6098926440472492

R2 Test Score: -1.1552672976079153e+22

Intercept: -2.909670123625102e+18

Coefficient: [1.91489567e+04 2.49117721e+04 8.14156176e+04 1.77261151e+03

1.90251959e+04 5.26717233e+05 -3.61645471e+05 1.80704109e+05


```

1.29117522e+05 1.23424043e+05 3.69725164e+05 9.34006371e+05
3.38667524e+05 2.31691087e+06 1.14969826e+05 1.93744193e+05
2.27014659e+04 5.38048172e+04 7.52325402e+05 2.90967012e+18
2.90967012e+18 2.90967012e+18 2.90967012e+18 2.90967012e+18
2.90967012e+18 2.90967012e+18 2.90967012e+18 2.90967012e+18
2.90967012e+18 9.78885051e+04]
MSE Train: 52518539058.11955
MSE Test: 1.568101542566779e+33
----- Q-Q Plot -----

```

<ipython-input-112-dbf471e5a3b>:24: UserWarning: Matplotlib is currently using module://ipykernel.pylab.backend_inline, which is a non-GUI backend, so cannot show the figure.

```
fig.show()
```



Some takeaways from the model

- R2 Stat got worse, we're still not really getting anywhere with the model
- P-values are now all at 0, which is good
- Mean squared error is the worst that it has ever been
- JB test showing that our data could not be further from normal if we tried
- F stat indicates that our data is not homoscedastic, and it's worse somehow
- Overall, we didn't get anywhere, but I have some hope trying this on our scaled data

```
In [127]: cv_scaled_df.dropna(inplace=True)
```

```
In [128]: X_scaled_cv = cv_scaled_df.drop('price_log', axis=1)
y_log_cv = cv_scaled_df['price_log']
```

```
In [129]: #Start up selector that will give us the 5 best features
lr = LinearRegression()
selector = RFE(lr, n_features_to_select = 5)
selector = selector.fit(X_scaled_cv, y_log_cv.ravel()) # convert y to 1d np
selector.support_
selected_columns = scaled_df.columns[selector.support_ ]
```

```
In [130]: X_scaled_cv.shape
```

```
Out[130]: (21589, 37)
```

```
In [131]: y_log_cv.shape
```

```
Out[131]: (21589,)
```

```
In [132]: linreg_process(X_scaled_cv[selected_columns], y_log_cv)
```

```

=====
                        OLS Regression Results
=====
=====
Dep. Variable:          price_log    R-squared:
0.004
Model:                  OLS    Adj. R-squared:
0.004
Method:                Least Squares    F-statistic:
13.63
Date:                  Wed, 03 Feb 2021    Prob (F-statistic):          2.6
7e-13
Time:                  20:27:32    Log-Likelihood:          -1
2598.
No. Observations:      16191    AIC:          2.52
1e+04
Df Residuals:          16185    BIC:          2.52
5e+04
Df Model:              5
Covariance Type:       nonrobust
=====
=====
                        coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const          13.0424      0.004    3149.224      0.000      13.034      1
3.050
grade_7         0.0197      0.007      2.854      0.004      0.006
0.033
grade_8         0.0349      0.007      5.297      0.000      0.022
0.048
grade_9         0.0324      0.006      5.792      0.000      0.021
0.043
grade_10        0.0288      0.005      5.929      0.000      0.019
0.038
grade_11        0.0165      0.005      3.587      0.000      0.007
0.026
=====
=====
Omnibus:          660.939    Durbin-Watson:
2.007
Prob(Omnibus):    0.000    Jarque-Bera (JB):          84
3.012
Skew:             0.439    Prob(JB):          8.76
e-184
Kurtosis:         3.691    Cond. No.
3.11
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


----- Scoring Stats -----

R2 Train Score: 0.004192454703602966

```
R2 Test Score: 0.0018149917275585015
Intercept: 13.04235344043171
Coefficient: [0.01965861 0.034929 0.03242326 0.02881512 0.01649134]
MSE Train: 0.2775519047375082
MSE Test: 0.27188522372178453
----- Q-Q Plot -----
```

```
<ipython-input-112-dbf471e5a3b>:24: UserWarning: Matplotlib is currently
using module://ipykernel.pylab.backend_inline, which is a non-GUI backen
d, so cannot show the figure.
```

```
fig.show()
```



Some takeaways from the model

- Obviously something terrible happened here...

Log-Transform X Variables w/ Scaler

This time we are going to log transform everything, then use the `StandardScaler()`, the biggest issue our data was having was it's lack of normality, so this is a last stand for getting that appropriate.

```
In [133]: data_preprocessed.head()
```

```
Out[133]:
```

| | price | sqft_living | sqft_above | sqft_living15 | bedrooms_2 | bedrooms_3 | bedrooms_4 | bedrooms |
|---|----------|-------------|------------|---------------|------------|------------|------------|----------|
| 0 | 221900.0 | 1180 | 1180 | 1340 | 0 | 1 | 0 | |
| 1 | 538000.0 | 2570 | 2170 | 1690 | 0 | 1 | 0 | |
| 2 | 180000.0 | 770 | 770 | 2720 | 1 | 0 | 0 | |
| 3 | 604000.0 | 1960 | 1050 | 1360 | 0 | 0 | 1 | |
| 4 | 510000.0 | 1680 | 1680 | 1800 | 0 | 1 | 0 | |

5 rows × 39 columns

```
In [134]: #Start with removing
data_preprocessed.drop(['price_log'], axis=1, inplace=True)
```

```
In [135]: continuous = ['price', 'sqft_living', 'sqft_above', 'sqft_living15']
data_cont = data_preprocessed[continuous]
log_names = [f'{column}_log' for column in data_cont.columns]
preped_log = np.log(data_preprocessed[continuous])
preped_log.columns = log_names
```

```
In [136]: data_preprocessed.head()
```

```
Out[136]:
```

| | price | sqft_living | sqft_above | sqft_living15 | bedrooms_2 | bedrooms_3 | bedrooms_4 | bedrooms |
|---|----------|-------------|------------|---------------|------------|------------|------------|----------|
| 0 | 221900.0 | 1180 | 1180 | 1340 | 0 | 1 | 0 | |
| 1 | 538000.0 | 2570 | 2170 | 1690 | 0 | 1 | 0 | |
| 2 | 180000.0 | 770 | 770 | 2720 | 1 | 0 | 0 | |
| 3 | 604000.0 | 1960 | 1050 | 1360 | 0 | 0 | 1 | |
| 4 | 510000.0 | 1680 | 1680 | 1800 | 0 | 1 | 0 | |

5 rows × 38 columns

```
In [137]: data_preprocessed.drop(continuous, axis=1, inplace=True)
```

```
In [138]: full_preped_df = pd.concat([prep_log, data_preprocessed],axis=1)
full_preped_df.head()
```

Out[138]:

| | price_log | sqft_living_log | sqft_above_log | sqft_living15_log | bedrooms_2 | bedrooms_3 | bedrooms_ |
|---|-----------|-----------------|----------------|-------------------|------------|------------|-----------|
| 0 | 12.309982 | 7.073270 | 7.073270 | 7.200425 | 0 | 1 | |
| 1 | 13.195614 | 7.851661 | 7.682482 | 7.432484 | 0 | 1 | |
| 2 | 12.100712 | 6.646391 | 6.646391 | 7.908387 | 1 | 0 | |
| 3 | 13.311329 | 7.580700 | 6.956545 | 7.215240 | 0 | 0 | |
| 4 | 13.142166 | 7.426549 | 7.426549 | 7.495542 | 0 | 1 | |

5 rows × 38 columns

```
In [139]: scaler = StandardScaler()
prep_X = full_preped_df.drop(['price_log'], axis=1)
y_log = full_preped_df['price_log']
X_scaled = scaler.fit_transform(prep_X)
```

```
In [140]: linreg_process(prepare_X, y_log)
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price_log    R-squared:
0.605
Model:                  OLS    Adj. R-squared:
0.604
Method:                Least Squares    F-statistic:
687.0
Date:                  Wed, 03 Feb 2021    Prob (F-statistic):
0.00
Time:                  20:27:33    Log-Likelihood:          -5
044.3
No. Observations:      16194    AIC:                  1.01
6e+04
Df Residuals:          16157    BIC:                  1.04
5e+04
Df Model:              36
Covariance Type:      nonrobust
=====
=====
                        coef      std err          t      P>|t|      [0.025
0.975]
-----
const                8.0974      0.105      77.057      0.000      7.891
8.303
sqft_living_log       0.3651      0.026     13.836      0.000      0.313
0.417
sqft_above_log       -0.0261      0.026     -1.000      0.317     -0.077
0.025
sqft_living15_log     0.2476      0.013     18.982      0.000      0.222
0.273
bedrooms_2            0.0003      0.028      0.011      0.991     -0.055
0.056
bedrooms_3           -0.1396      0.028     -4.910      0.000     -0.195
-0.084
bedrooms_4           -0.1524      0.029     -5.214      0.000     -0.210
-0.095
bedrooms_5           -0.1229      0.031     -3.976      0.000     -0.183
-0.062
bedrooms_6           -0.1149      0.037     -3.076      0.002     -0.188
-0.042
bedrooms_7           -0.0902      0.067     -1.346      0.178     -0.222
0.041
bedrooms_8            0.0935      0.110      0.849      0.396     -0.122
0.309
bedrooms_9           -0.0392      0.175     -0.224      0.823     -0.382
0.304
bedrooms_10           0.0582      0.236      0.246      0.806     -0.405
0.522
bedrooms_11           0.0315      0.332      0.095      0.925     -0.620
0.683
bathrooms_2.0        -0.0357      0.009     -3.844      0.000     -0.054
-0.017

```


| | | | | | |
|-------------------------|---------|-------|--------|-------|--------|
| bathrooms_3.0 -0.059 | -0.0825 | 0.012 | -6.991 | 0.000 | -0.106 |
| bathrooms_4.0 0.052 | 0.0172 | 0.018 | 0.953 | 0.340 | -0.018 |
| bathrooms_5.0 0.182 | 0.1102 | 0.036 | 3.027 | 0.002 | 0.039 |
| bathrooms_6.0 0.406 | 0.2278 | 0.091 | 2.510 | 0.012 | 0.050 |
| bathrooms_7.0 0.065 | -0.3155 | 0.194 | -1.623 | 0.105 | -0.696 |
| bathrooms_8.0 0.527 | 0.1618 | 0.186 | 0.868 | 0.385 | -0.203 |
| floors_1.5 0.240 | 0.2206 | 0.010 | 22.556 | 0.000 | 0.201 |
| floors_2.0 0.019 | 0.0026 | 0.008 | 0.312 | 0.755 | -0.014 |
| floors_2.5 0.255 | 0.1948 | 0.031 | 6.328 | 0.000 | 0.134 |
| floors_3.0 0.214 | 0.1799 | 0.017 | 10.420 | 0.000 | 0.146 |
| floors_3.5 0.507 | 0.2593 | 0.126 | 2.051 | 0.040 | 0.011 |
| waterfront_1.0 0.638 | 0.5727 | 0.033 | 17.299 | 0.000 | 0.508 |
| grade_4 0.287 | 0.1543 | 0.068 | 2.277 | 0.023 | 0.021 |
| grade_5 0.174 | 0.1213 | 0.027 | 4.533 | 0.000 | 0.069 |
| grade_6 0.333 | 0.2985 | 0.018 | 16.890 | 0.000 | 0.264 |
| grade_7 0.506 | 0.4730 | 0.017 | 28.356 | 0.000 | 0.440 |
| grade_8 0.692 | 0.6567 | 0.018 | 36.660 | 0.000 | 0.622 |
| grade_9 0.928 | 0.8886 | 0.020 | 44.152 | 0.000 | 0.849 |
| grade_10 1.117 | 1.0727 | 0.023 | 47.069 | 0.000 | 1.028 |
| grade_11 1.298 | 1.2422 | 0.029 | 43.324 | 0.000 | 1.186 |
| grade_12 1.510 | 1.4237 | 0.044 | 32.339 | 0.000 | 1.337 |
| grade_13 1.989 | 1.7665 | 0.114 | 15.558 | 0.000 | 1.544 |
| basement_1 0.143 | 0.1223 | 0.011 | 11.433 | 0.000 | 0.101 |

=====

=====

| | | | |
|----------------|-------|-------------------|-----|
| Omnibus: | 1.364 | Durbin-Watson: | |
| 1.997 | | | |
| Prob(Omnibus): | 0.505 | Jarque-Bera (JB): | |
| 1.336 | | | |
| Skew: | 0.016 | Prob(JB): | |
| 0.513 | | | |
| Kurtosis: | 3.030 | Cond. No. | 5.6 |
| 0e+15 | | | |

=====

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 8.83e-26. This might indicate that there are

strong multicollinearity problems or that the design matrix is singular.

----- Scoring Stats -----

R2 Train Score: 0.6048743096944382

R2 Test Score: 0.6152590685646006

Intercept: 8.907141974483777

Coefficient: [3.65056551e-01 -2.60842552e-02 2.47647475e-01 3.12192417e-04

-1.39612358e-01 -1.52369775e-01 -1.22905173e-01 -1.14869696e-01

-9.01963856e-02 9.34857491e-02 -3.92274948e-02 5.82292576e-02

3.14879045e-02 -3.56516479e-02 -8.24980751e-02 1.71638668e-02

1.10221467e-01 2.27816468e-01 -3.15515520e-01 1.61807567e-01

2.20598839e-01 2.56324544e-03 1.94807808e-01 1.79867925e-01

2.59317701e-01 5.72700532e-01 -6.55427380e-01 -6.88415470e-01

-5.11249119e-01 -3.36695510e-01 -1.53054441e-01 7.88312043e-02

2.62915003e-01 4.32412031e-01 6.13962772e-01 9.56720910e-01

1.22299784e-01]


MSE Train: 0.10916527927993505

MSE Test: 0.107635270143443

----- Q-Q Plot -----

<ipython-input-112-dbf471e5a3b>:24: UserWarning: Matplotlib is currently using module://ipykernel.pylab.backend_inline, which is a non-GUI backend, so cannot show the figure.

fig.show()



In [141]: cross_val(prepare_X, y_log)

0.10949017844190456

Some takeaways

- R2 Score:
 - Ultimately despite our best efforts at scaling the data, the original model managed to perform the best for the percentage of variance explained by the model, but this effort at right around 61% isn't awful.
 - Test R2 was about 1.5% better than the train R2 indicating some mild over-fitting, but nothing out of the ordinary
- P-values:
 - We had 13 variables above the alpha level, meaning that there was a lot of reason to believe that many of these were not having an effect on our target variable
- Mean Squared Error (MSE):
 - This is where we saw the greatest effect on our model, we did not have a whole lot of errors and overall our regression line was pretty close to the actual values
 - It wasn't 0 (indicating an overfit perfect line), but there wasn't any indication of overfitting when comparing the train and test data
- JB Test:
 - We also performed very well here and showed that we were able to get very normal data out of this set, especially since it started so far off
 - This was also absolutely confirmed in our Q-Q plot
- Overall:
 - This wasn't a terrible model, and essentially can account for roughly 60% of an accurate home price (in layman's terms)
 - We didn't have any significant overfitting, and were able to get very normal data, there is more work that needs to be done to achieve a more accurate model, but we can get into that in the conclusion

Conclusions

I think this is a great starting point to assist with someone in setting a selling point for their home. with this current model I was able to predict roughly 60% accurate home prices, which leaves a lot to be desired for the homeseller, but there is more work that could certainly be done to improve this further. There is a lot that goes into home selling and unfortunately, due to time constraints, we were not able to model with one of the most important sets of variables...location. I think that a lot of our errors in the modeling came as a result of not having location data and you could see that in some of the outliers. A small home in the Downtown Seattle will be far more expensive than a large home in the furthest reaches of King County. In some ways the issue with this data set was that it was too large, and some future work would be breaking it down by zipcode or geographic area (using latitude and longitude).

Further Work

- As I stated in the conclusions section, I think the most important aspect would be breaking down the zipcode and/or longitude and latitude to make smaller models based on more relevant data
- Another thing I would like to do is to break down and separate some of the larger homes and larger lots

-

- Many of these are likely on large plots of cheaper land, and really shouldn't be compared with homes in more desirable locations